



# Optimizing gpucc (part 2)

Jingyue Wu, Artem Belevich, Eli Bendersky, Mark Heffernan, Chris Leary,  
**Jacques Pienaar**, Bjarke Rouné, Rob Springer, Xuetian Weng, Robert Hundt

# Example

"I only need the high-bits of a multiplication"

# Where does this arise?

- Input program\*

```
__global__ void divide_by_5(unsigned int* x, unsigned int* y) {  
    int i = blockIdx.x * blockDim.x + threadIdx.x;  
    y[i] = x[i] / 5;  
}
```

- Generated PTX:

```
mul.lo.s64  %rd8, x[i], 3435973837;  
shr.u64     %rd9, %rd8, 34;  
st.u32      y[i], %rd9;
```

# Why?

- Multiply with magic number instead of dividing by integer
  - From Hacker's Delight

$$\begin{aligned}\left\lfloor \frac{n}{5} \right\rfloor &= \left\lfloor \frac{n}{5} \times \frac{2^k}{2^k} \right\rfloor \\ &= \left\lfloor m \times \frac{n}{2^k} \right\rfloor \\ &= \left\lfloor \frac{2^k + e}{5} \times \frac{n}{2^k} \right\rfloor \\ &= \left\lfloor \frac{n}{5} + \frac{e}{5} \times \frac{n}{2^k} \right\rfloor\end{aligned}$$

# Why?

- Multiply with magic number instead of dividing by integer
  - From Hacker's Delight

$$\left\lfloor \frac{n}{5} \right\rfloor = \left\lfloor \frac{n}{5} + \frac{e}{5} \times \frac{n}{2^k} \right\rfloor$$

$$k = 32 + 2$$

$$e = 5 - 2^k \mod 5 = 1$$

$$\frac{e}{5} \times \frac{n}{2^k} = \frac{1}{5} \times \frac{n}{2^k} < \frac{1}{5} \times \frac{2^{32}}{2^{34}} < \frac{1}{5}$$

$$m = \left\lceil \frac{2^k}{5} \right\rceil = 3435973837$$

## But why 64-bit mul?

- Multiply by 64-bits and shift right by 34
- Why not: mul hi & shift by 2?

### PTX ISA

```
// extended-precision multiply: [r3,r2,r1,r0] = [r5,r4] * [r7,r6]
[ ... ]
mul.hi.u32 r1,r4,r6;           // r1=(r4*r6).[63:32], no carry-out
[ ... ]
```

# Finding out what the backend is doing

1. Compile with -v (lots of output)

```
$ ./bin/clang++ div.cu -o div -L<SDK>/lib64  
--cuda-gpu-arch=sm_35 -lcudart_static -  
lcuda -ldl -lrt -pthread -v
```

2. Compile for device only

```
$ ./bin/clang++ div.cu -o div.o --cuda-gpu-  
arch=sm_35 --cuda-device-only -c
```

# Finding out what the backend is doing

```
$ <cmd> -mllvm -help-hidden | grep isel
```

<code>-print-isel-input</code>	- Print LLVM IR input to isel pass
<code>-view-isel-dags</code>	- Pop up a window to show isel dags as they are selected

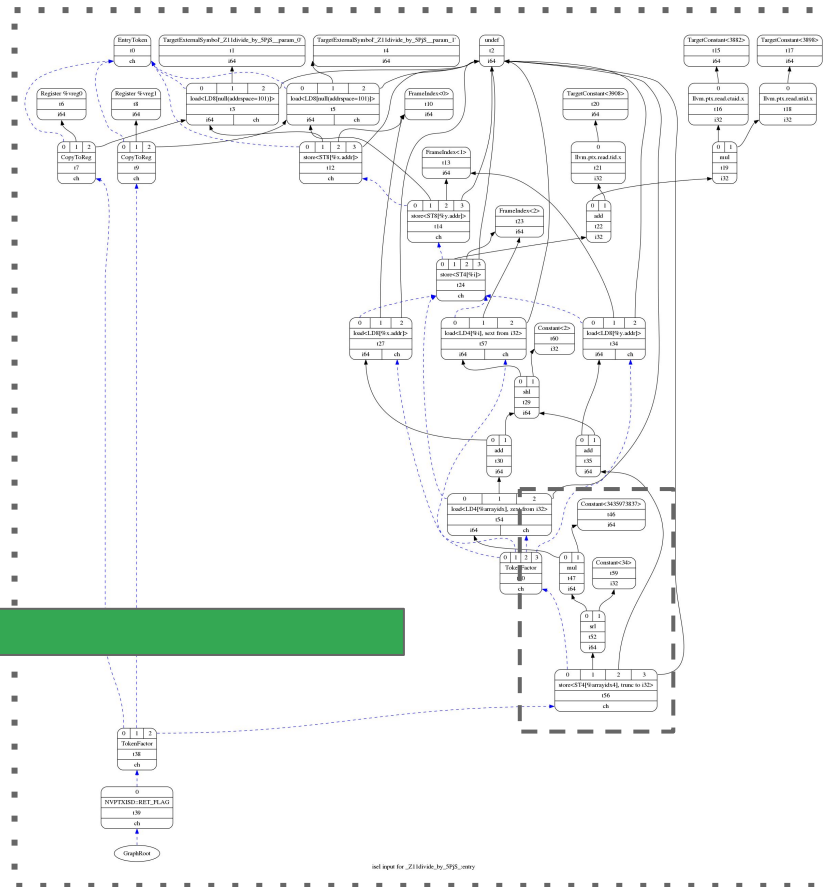


## Print ISEL input

```
__global__ void divide_by_5(  
    unsigned int* x, unsigned int* y) {  
  
    int i = blockIdx.x * blockDim.x +  
           threadIdx.x;  
  
    y[i] = x[i] / 5;  
}
```

```
define void @_Z11divide_by_5PjS_(i32* %x, i32* %y) #2 {  
entry:  
    %x.addr = alloca i32*, align 8  
    %y.addr = alloca i32*, align 8  
    %i = alloca i32, align 4  
    store i32* %x, i32** %x.addr, align 8  
    store i32* %y, i32** %y.addr, align 8  
    %0 = call i32 @llvm.ptx.read.ctaid.x() #4  
    %1 = call i32 @llvm.ptx.read.ntid.x() #4  
    %mul = mul i32 %0, %1  
    %2 = call i32 @llvm.ptx.read.tid.x() #4  
    %add = add i32 %mul, %2  
    store i32 %add, i32* %i, align 4  
    %3 = load i32, i32* %i, align 4  
    %idxprom = sext i32 %3 to i64  
    %4 = load i32*, i32** %x.addr, align 8  
    %arrayidx = getelementptr inbounds i32, i32* %4, i64 %idxprom  
    %5 = load i32, i32* %arrayidx, align 4  
    %div = udiv i32 %5, 5  
    %6 = load i32, i32* %i, align 4  
    %idxprom3 = sext i32 %6 to i64  
    %7 = load i32*, i32** %y.addr, align 8  
    %arrayidx4 = getelementptr inbounds i32, i32* %7, i64 %idxprom3  
    store i32 %div, i32* %arrayidx4, align 4  
    ret void  
}
```

Google



# NVPTX backend

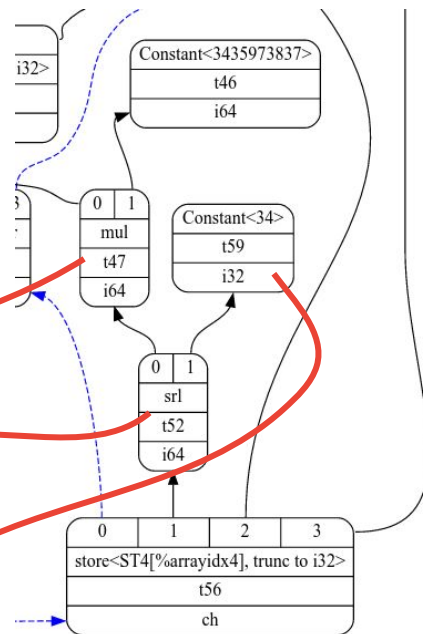
- NVPTX : PTX backend originally contributed by NVIDIA
- NVPTXInstrInfo.td : TableGen file with instruction information

```
defm SRL      : RSHIFT_FORMAT<"shr.u", srl>;
defm MULHS    : I3<"mul.hi.s", mulhs>;
multiclass I3<string OpcStr, SDNode OpNode> {
  [...]
  def i32ri : NVPTXInst<(outs Int32Regs:$dst), (ins Int32Regs:$a, i32imm:$b),
    !strconcat(OpcStr, "32 \t$dst, $a, $b;"),
    [(set Int32Regs:$dst, (OpNode Int32Regs:$a, imm:$b))]>;
  [...]
}
```

# Aha, add special matcher

- No magic: look and match
- One (very small) peephole optimization

```
def : Pat<(srl (mul Int64Regs:$a, (i64 imm:$b)), (i32 34))  
  (SRLi32ri (MULTHUi32ri Int32Regs:$a, imm:$b), 2)>;
```



## Not so fast ...

```
Anonymous_516: (SRLi32ri:i32 (MULTHUi32ri:i32 Int32Regs:
<empty>:$a, (imm:i32):$b), 2:i32)
Included from cgo-tut/llvm/lib/Target/NVPTX/NVPTX.td:19:
cgo-tut/llvm/lib/Target/NVPTX/NVPTXInstrInfo.td:1091:1: error: In
anonymous_516: Type inference contradiction found, merging 'i64'
into 'i32'
def : Pat<(srl (mul Int64Regs:$a, (i64 imm:$b)), (i32 34)),
^

def : Pat<(srl (mul Int64Regs:$a, (i64 imm:$b)), (i32 34)),
      (SRLi32ri (MULTHUi32ri Int32Regs:$a, imm:$b), 2)>;
```

# Debug output to the rescue

\$ <cmd> -mllvm -debug -mllvm -print-after-all

- Find the last udiv:

```
Combining: t33: i32 = udiv t31, Constant:i32<5>
```

```
... into: t44: i32 = srl t42, Constant:i32<2>
```

```
[...]
```

```
Combining: t42: i32 = mulhu t31, Constant:i32<-858993459>
```

```
... into: t50: i32 = truncate t49
```

3435973837

# Trace back change

- Search through LLVM files for "Combining:"
  - Leads to "lib/CodeGen/SelectionDAG/DAGCombiner.cpp"
- Change happens in mulhu, so look at visitMULHU:

```
SDValue DAGCombiner::visitMULHU(SDNode *N) {  
    [...]  
    // If the type twice as wide is legal, transform the mulhu to a wider multiply  
    // plus a shift.  
    if (VT.isSimple() && !VT.isVector()) {  
        [...]  
    }  
  
    return SDValue();  
}
```

# Trace back change

- Search through LLVM files for "Combining:"
  - Leads to "lib/CodeGen/SelectionDAG/DAGCombiner.cpp"
- Change happens in mulhu, so look at visitMULHU:

```
SDValue DAGCombiner::visitMULHU(SDNode *N) {  
    [...]  
    // If the type twice as wide is legal and not targeting NVPTX, transform the  
    // mulhu to a wider multiply plus a shift.  
    Triple TT = DAG.getTarget().getTargetTriple();  
    if (!TT.isNVPTX() && VT.isSimple() && !VT.isVector()) {  
        [...]  
    }  
  
    return SDValue();  
}
```



# Difference in resultant output

old.s

```
+-- 25 lines: // Generated by LLVM NVPTX Back-End
.param .u64 _Z11divide_by_5PjS__param_1
```

```
)
{
```

```
.local .align 8 .b8 __local_depot1[24];
.reg .b64 %SP;
.reg .b64 %SPL;
.reg .s32 %r<6>;
.reg .s64 %rd<12>;
```

```
+-- 15 lines
```

```
add.s64 %rd6, %rd4, %rd5;
ld.u32 %rd7, [%rd6];
mul.lo.s64 %rd8, %rd7, 3435973837;
shr.u64 %rd9, %rd8, 34;
ld.u64 %rd10, [%SP+8];
add.s64 %rd11, %rd10, %rd5;
st.u32 [%rd11], %rd9;
ret;
```

```
}
```

new.s

```
+-- 25 lines: // Generated by LLVM NVPTX Back-End
.param .u64 _Z11divide_by_5PjS__param_1
```

```
)
{
```

```
.local .align 8 .b8 __local_depot1[24];
.reg .b64 %SP;
.reg .b64 %SPL;
.reg .s32 %r<9>;
.reg .s64 %rd<9>;
```

```
+-- 15 lines
```

```
add.s64 %rd6, %rd4, %rd5;
ld.u32 %r6, [%rd6];
mul.hi.u32 %r7, %r6, -858993459;
shr.u32 %r8, %r7, 2;
ld.u64 %rd7, [%SP+8];
add.s64 %rd8, %rd7, %rd5;
st.u32 [%rd8], %r8;
ret;
```

```
}
```


## But is it needed?

- Type widening for multhu introduced in <http://llvm.org/viewvc/llvm-project?view=revision&revision=121696>  
(Mon Dec 13 00:39:01 PST 2010)

### **Intel Optimization Reference Manual, Assembly/Compiler Coding Rule 21:**

Favor generating code using imm8 or imm32 values instead of imm16 values. If imm16 is needed, load equivalent imm32 into a register and use the word value in the register instead.

# Does it help? (SHOC)



**"Division is really slow! But I know some of my divisors are powers of two ..."**


# Example input

- Input kernel

```
__global__ void divide_by_a(int a, unsigned int* x, unsigned int* y) {  
    int i = blockIdx.x * blockDim.x + threadIdx.x;  
    y[i] = x[i] / a;  
}
```

- 'a' is not constant so regular power of 2 transformation won't apply
- BUT division is so much slower than compare ...

# What should the pass do?



# Fast path for avoiding divides at runtime

- Common request == utility exists!
- Existing transform utility to bypass slow division:

## **BypassSlowDivision.cpp**

```
// This file contains an optimization for div and rem on architectures that  
// execute short instructions significantly faster than longer instructions.  
// For example, on Intel Atom 32-bit divides are slow enough that during  
// runtime it is profitable to check the value of the operands, and if they are  
// positive and less than 256 use an unsigned 8-bit divide.
```

- For GPUs it might pay to be even more aggressive!
  - Or not - profiling your app will indicate!
  - Is pass needed?

# Add a new pass

```
struct NVPTXBypassDivPowerTwo : public FunctionPass {  
    static char ID;  
  
    NVPTXBypassDivPowerTwo() : FunctionPass(ID) {}  
  
    bool runOnFunction(Function &F) override;  
  
    const char *getPassName() const override {  
        return "Bypass divide power two";  
    }  
};
```


```
void NVPTXPassConfig::addCodeGenPrepare() {  
    addPass(createBypassDivPowerTwo());  
    TargetPassConfig::addCodeGenPrepare();  
}
```

```
namespace llvm {  
    void initializeNVPTXBypassDivPowerTwoPass  
        (PassRegistry&);  
}  
  
INITIALIZE_PASS(NVPTXBypassDivPowerTwo,  
    "Nvptx-bypass-div-power-two",  
    "Insert runtime check to use shift instead of div",  
    false, false)  
  
FunctionPass *llvm::createBypassDivPowerTwo() {  
    return new NVPTXBypassDivPowerTwo();  
}
```



# Working of the pass

- `runOnFunction(Function &F)`
  - For every basic block
    - For every instruction in the basic block
      - Is this a binary operation that performs division?
        - a. Split original block in 2 (before divide, after divide)
        - b. Add 2 new basic blocks (slow path, fast path)
        - c. Insert check for power of 2 (use population counting)
        - d. Fast path: shift input trailing-zeros (divisor) left
        - e. Slow path: divide input by divisor
        - f. After divide: join via phi-node



# Difference in resultant output

old.s

```
+-- .param .u64 _Z11divide_by_aiPjS__param_2) {  
ld.u32 %r71, [%rd6];  
ld.u32 %r82, [%SP+0];  
  
div.u32 %r94, %r71, %r82;  
  
ld.u64 %rd8, [%SP+16];  
  
add.s64 %rd10, %rd8, %rd5;  
st.u32 [%rd10], %r94;  
ret;  
}
```

new.s

```
+-- .param .u64 _Z11divide_by_aiPjS__param_2) {  
ld.u32 %r71, [%rd6];  
ld.u32 %r82, [%SP+0];  
popc.b32 %r12, %r82;  
setp.ne.s32 %p1, %r12, 1;  
@%p1 bra LBB2_2;  
bra.uni LBB2_1;  
LBB2_1:  
clz.b32 %r13, %r2;  
mov.u32 %r14, 31;  
sub.s32 %r15, %r14, %r13;  
shr.u32 %r3, %r71, %r15;  
mov.u32 %r16, %r3;  
bra.uni LBB2_3;  
LBB2_2:  
div.u32 %r94, %r71, %r82;  
mov.u32 %r16, %r94;  
bra.uni LBB2_3;  
LBB2_3:  
mov.u32 %r5, %r16;  
ld.s32 %rd7, [%SP+24];  
ld.u64 %rd8, [%SP+16];  
shl.b64 %rd9, %rd7, 2;  
add.s64 %rd10, %rd8, %rd9;  
st.u32 [%rd10], %r5;  
ret;  
}
```



# Contributing to gpucc

**Jingyue Wu**, Artem Belevich, Eli Bendersky, Mark Heffernan, Chris Leary,  
**Jacques Pienaar**, Bjarke Rouné, Rob Springer, Xuetian Weng, Robert Hundt

**"We have a performant open-source GPU compiler. Now what?"**

# Compiler development

- Compilers are tuned to provide good performance on important\* applications and benchmarks
- Where important can be 1) benchmarks, 2) widely used kernels, or 3) applications that the compiler developer cares about
- gpuc performs better than nvcc on our internal applications as that was our first target
  - It performs well on open-source benchmarks too!

**But it should be the best for applications you care about too!**

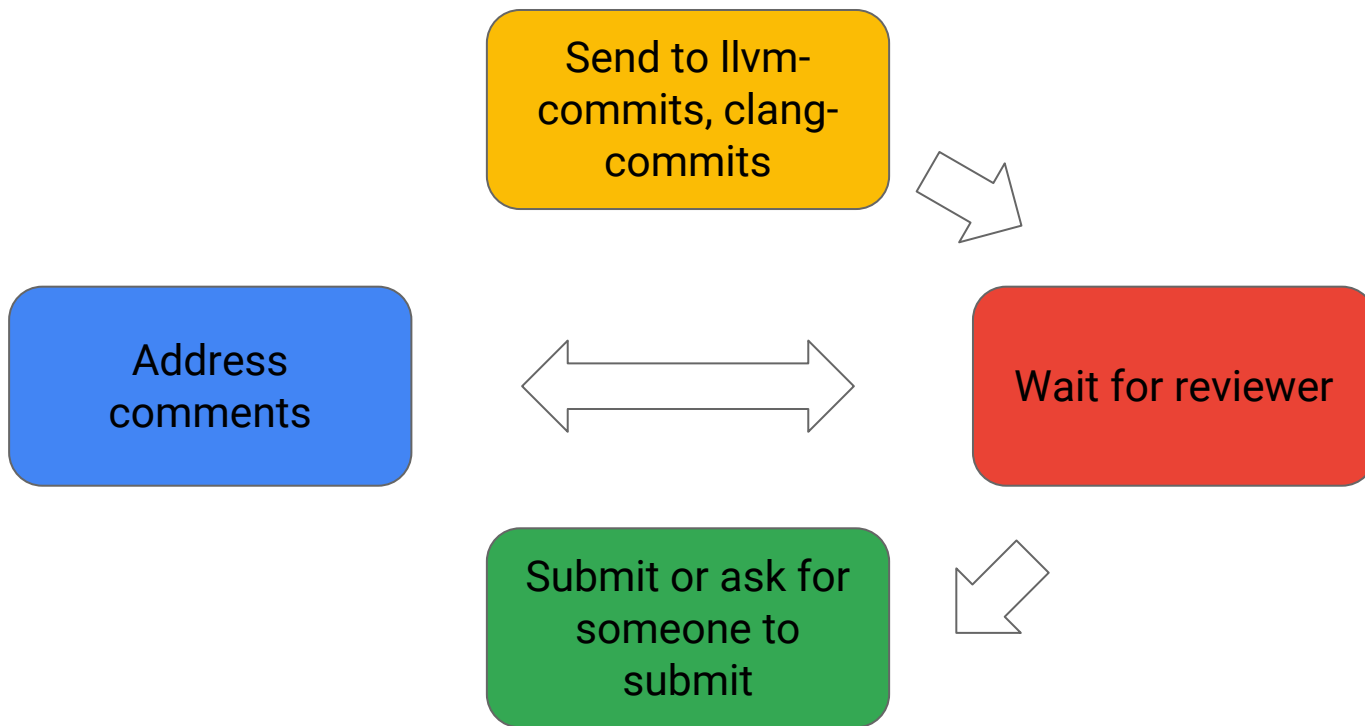
# How to contribute to LLVM

- License
  - Published under the "University of Illinois/NCSA Open Source License" BSD-style
  - Some of them are also published under MIT licence (dual-license)
  - No copyright assignment
- Community
  - Friendly
  - Mainly professional (Apple, Google, ARM, Linaro, Intel, etc)
  - But many individual and academic contributors
  - Usually fast to answer to comments/questions
- Resources
  - [LLVM Developer Policy](#)
  - ["How to contribute to LLVM, Clang, Etc"](#) talk by Sylvestre Ledru @ FOSDEM 2014

# The LLVM/Clang community

	LLVM	Clang
Project Activity	Very High	Very High
Homepage	<a href="http://llvm.org/">http://llvm.org/</a>	<a href="http://clang.llvm.org/">http://clang.llvm.org/</a>
Project License	NCSA	NCSA
Contributors (All Time)	613 developers	459 developers
Commits (All Time)	134,904 commits	62,633 commits
Initial Commit	over 14 years ago	over 8 years ago
Contributors (Past 12 Months)	336 developers	265 developers
Commits (Past 12 Months)	15,426 commits	5,983 commits

# Contributing a patch





# Getting help

- Mailing lists:
  - LLVM
    - LLVM-dev
    - LLVM-commits
  - Clang
    - cfe-users
    - cfe-dev
    - cfe-commits
- IRC: [irc.oftc.net](https://irc.oftc.net) #llvm

Keep patches small and focussed

Discuss the design on the mailing list, before starting the development

Review patches

# Organization of LLVM source code

- `$llvm/lib/Target/NVPTX`  
NVPTX backend & NVPTX specific optimizations
- `$llvm/lib/Transforms/Utils`  
General transformation utils
- `$llvm/CodeGen/` and `$llvm/CodeGen/SelectionDAG`  
General code generations and transformations on instruction selection DAG
- `$llvm/tools/clang/lib/Sema/SemaCUDA.cpp`  
Semantic analysis for CUDA constructs
- `$llvm/tools/clang/include/clang/Basic/BuiltinsNVPTX.def`,  
`$llvm/tools/clang/lib/CodeGen/CGCUDABuiltin.cpp`  
CUDA specific builtins

# Using out-of-tree

- Use gpucc/Clang/LLVM to create tools & plugins
- Useful resources:
  - "Writing an LLVM Pass"  
<http://llvm.org/docs/WritingAnLLVMPass.html>
  - "Tutorial: Building, Testing and Debugging a Simple out-of-tree LLVM Pass"  
<http://www.llvm.org/devmtg/2015-10/#tutorial1>
  - "Samples for using LLVM and Clang as a library"  
<http://eli.thegreenplace.net/2014/samples-for-using-llvm-and-clang-as-a-library/>  
[https://github.com/eliben/llvm-clang-samples/tree/master/src\\_clang](https://github.com/eliben/llvm-clang-samples/tree/master/src_clang)  
<http://eli.thegreenplace.net/tag/llvm-clang>
  - "LLVM for Grad Students"  
<http://adriansampson.net/blog/llvm.html>

# Action Items

# Concrete bugs and feature requests

- Frontend support
  - Texture memory
  - C++14
- NVPTX Backend
  - More intrinsics
- CUDA runtime support
  - dynamic allocation and deallocation
- Micro-optimizations / peephole

# Loops with strided accesses

```
__global__ void kernel(float* output, int numIterations) {  
    for (int i = 0; i < numIterations; i += 2) {  
        output[i] = output[i+1];  
    }  
}
```

- Good: loop gets unrolled a few times
- Bad:
  - gpucc fails to fold the offsets from the induction variable into the constant addressing mode
  - Uses 64-bit indexing even where 32-bit would suffice

# Automatically generated memset sub-optimal

```
__global__ void kernel(float* output, int N) {  
    for (int i = 0; i < N; ++i) {  
        output[i] = 0;  
    }  
}
```

- The PTX output by gcudacc changes that to be 1 byte at a time
  - Idiom recognize as memset
  - Generated memset bad

# Unnecessary copies

- Why copy into r16 before copying into r5?

- Caveat: PTXAS is an optimizing assembler. Performs more optimizations than one might expect.

new.s

```
.param .u64 _Z11divide_by_aiPjS__param_2) {  
+-- not important lines  
ld.u32 %r71, [%rd6];  
ld.u32 %r82, [%SP+0];  
popc.b32 %r12, %r82;  
setp.ne.s32 %p1, %r12, 1;  
@%p1 bra LBB2_2;  
bra.uni LBB2_1;  
LBB2_1:  
clz.b32 %r13, %r2;  
mov.u32 %r14, 31;  
sub.s32 %r15, %r14, %r13;  
shr.u32 %r3, %r71, %r15;  
mov.u32 %r16, %r3;  
bra.uni LBB2_3;  
LBB2_2:  
div.u32 %r94, %r71, %r82;  
mov.u32 %r16, %r94;  
bra.uni LBB2_3;  
LBB2_3:  
mov.u32 %r5, %r16;  
ld.s32 %rd7, [%SP+24];  
ld.u64 %rd8, [%SP+16];  
shl.b64 %rd9, %rd7, 2;  
add.s64 %rd10, %rd8, %rd9;  
st.u32 [%rd10], %r5;  
ret;  
}
```



# Research opportunities

## Past

- nvcc: CUDA  $\rightarrow$  *blackbox*  $\rightarrow$  PTX
- libnvvm: closed-source and works on outdated IR
- NVPTX codegen with little optimization

Present: complete open-source pipeline from CUDA to PTX


- Frontend: more language extensions? Other languages?
- Backend: more architectures (e.g., CUDA on AMDGPU)? Target SASS?
- **Optimizer: more optimizations?**

# Discover optimization opportunities

What optimizations do you currently have to perform in source code?


- Global memory coalescing
- Tune occupancy
- Avoid divergence
- Optimizations across host-device boundary
- Fast math
- More in *CUDA C Best practices* (<http://docs.nvidia.com/cuda/cuda-c-best-practices-guide/>)

# Global memory coalescing




```
void copy(float* input, float* output, int n) {  
    for (int i = 0; i < n; ++i)  
        output[i] = input[i];  
}
```

# Global memory coalescing




```
__global__ void copy(float* input, float* output, int n) {  
    for (int i = threadIdx.x * (n / 32);  
        i < (threadIdx.x + 1) * (n / 32); ++i) {  
        output[i] = input[i];  
    }  
}
```

# Uncoalesced access



```
__global__ void copy(float* input, float* output, int n) {  
    for (int i = threadIdx.x * (n / 32);  
        i < (threadIdx.x + 1) * (n / 32); ++i) {  
        output[i] = input[i];  
    }  
}
```

# Uncoalesced access



```
__global__ void copy(float* input, float* output, int n) {  
    for (int i = threadIdx.x * (n / 32);  
        i < (threadIdx.x + 1) * (n / 32); ++i) {  
        output[i] = input[i];  
    }  
}
```

# Coalesced access



```
__global__ void copy(float* input, float* output, int n) {  
    for (int block = 0; block < n / 32; ++block) {  
        int i = block * 32 + threadIdx.x;  
        output[i] = input[i];  
    }  
}
```


# Occupancy

Factors contributing to occupancy ([bit.ly/cuda-occupancy-calculator](https://bit.ly/cuda-occupancy-calculator))

- Number of registers each block uses
- Shared memory size
- Block size

Trade off between occupancy and latency

Feedback directed optimizations (FDO)?






# Divergence



```
if (threadIdx.x % 2 == 0) {
```

```
    foo();
```

A horizontal bar composed of 16 vertical segments, all colored blue, representing the threads that are executing the 'foo()' function.

```
} else {
```

```
    bar();
```

A horizontal bar composed of 16 vertical segments, all colored orange, representing the threads that are executing the 'bar()' function.


```
}
```




CPU:  $\max(\text{foo}, \text{bar})$

GPU:  $\text{foo} + \text{bar}$

# How divergence affects jump threading



`foo + bar + baz`



`foo + bar + baz * 2`

Divergence Analysis:  
<http://bit.ly/divergence-analysis>

# Optimizations across cross and device boundaries

```
__global__ void kernel(float* input,
                       float* output) {
    ... blockDim.x ...
}

float* input = cudaMemAlloc(...);
float* output = cudaMemAlloc(...);
kernel<<<grid_dim, 128>>>(input, output);
```

# Optimizations across cross and device boundaries

```
__global__ void kernel(float* input,
                       float* output) {
    ... blockDim.x ...
}

float* input = cudaMemAlloc(...);
float* output = cudaMemAlloc(...);
kernel<<<grid_dim, 128>>>(input, output);
```

Constant propagation

# Optimizations across cross and device boundaries

```
__global__ void kernel(float* input,  
                      float* output) {  
    ... blockDim.x ...  
}  
  
float* input = cudaMemAlloc(...);  
float* output = cudaMemAlloc(...);  
kernel<<<grid_dim, 128>>>(input, output);
```

Constant propagation

Alias analysis

# Current workaround: templates and annotations

```
__global__ void kernel(float* input,
                      float* output) {
    ... blockDim.x ...
}

float* input = cudaMemAlloc(...);
float* output = cudaMemAlloc(...);
kernel<<<grid_dim, 128>>>(input, output);
```



```
template <int kBlockDim>
__global__ void kernel(
    float* __restricted__ input,
    float* __restricted__ output) {
    ... kBlockDim ...
}

float* input = cudaMemAlloc(...);
float* output = cudaMemAlloc(...);
kernel<128><<<grid_dim, 128>>>(input, output);
```

Constant propagation  
Alias analysis

# Fast math: div → mul

```
// n = 1024
__global__ void foo(float *input, float *output, int n) {
    for (int i = 0; i < n; ++i)
        output[i] = input[i] / 255.0f;
}
```

	w/o fast math	w/ fast math
PTX	div.rn.f32	div.approx.ftz.f32 <b>input[i] * (1.0f / 255.0f)</b>
run time	990.72 us (2.6x)	373.92 us

# Takeaways from this tutorial

- The missions of gpucc
  - state-of-the-art platform for GPU compiler research
  - enables more industrial deployment
- In a good shape but can be improved a lot
- **Use and contribute to gpucc!** (<http://bit.ly/llvm-cuda>)